

P. Phani Hrushik Reddy

phanihrushik.10@gmail.com ❖ +91 863-954-8895 ❖ Hyderabad, India

Github: [@hrushik98](#) | LinkedIn: [Hrushik Pabbathireddy](#) | Leetcode: [phanihrushik](#)

WORK EXPERIENCE

Techolution

May 2025 – Present

AI Engineer, Fulltime, In Office

Hyderabad, India

- **Architected and scaled backend infrastructure** for AI-powered client solutions, enabling reliable deployment of production LLM services and supporting **20K+ inference requests/day** across multiple applications.
- Built and deployed internal AI-powered enterprise applications using **Angular (frontend)** and **Spring Boot (backend)** for **HarperCollins**, enabling seamless interaction with LLM-based services across multiple business workflows.
- **Fine-tuned, optimized, and served open-source LLMs** including **LLaMA-3**, **DeepSeek**, and **Qwen-VL**, improving response accuracy by **~35%** and reducing inference latency by **~40%** through quantization and optimized serving pipelines.
- **Built real-time speech AI pipelines** using **Whisper** and **Kokoro**, enabling automated transcription and natural voice synthesis with **sub-2s processing latency** for conversational AI workflows.
- **Designed and deployed scalable AI infrastructure on GCP**, orchestrating distributed model workloads capable of handling **10K+ concurrent inference jobs** with high availability.
- **Led rapid experimentation across ML and deep learning architectures**, accelerating internal prototyping cycles by **3×** and contributing to **5+ AI feature prototypes** adopted in client and internal solutions.

Multeway

Apr. 2024 – Nov. 2024

Fullstack Developer, Internship, In Office

Hyderabad, India

- Fine-tuned **LLaVA-1.5-7B (Vision LLM)** to detect illicit substance listings using multimodal inputs (image + text), improving moderation accuracy by **~90%** by training on contextual double-entendre patterns and visual cues.
- Developed and deployed the company's frontend web application using **React**, creating responsive UI components integrated with AI-driven backend services.
- Fine-tuned and quantized **LLaMA-2 7B** for chatbot query handling and optimized inference for low-latency responses; additionally fine-tuned domain-specific LLMs for customer support using RAG grounded on internal company knowledge.
- Designed and built the complete **FastAPI backend** deployed via Docker on **Kubernetes** (Helm-based provisioning), integrating **Kafka** for mobile notifications and **PGVector** for semantic + structured retrieval.
- Architected asynchronous pipelines using **RabbitMQ** + **Celery** for face verification workflows, processing queued image tasks with a TensorFlow-based fine-tuned face verification model.
- Implemented observability using **Prometheus** + **Grafana**, enabling real-time system monitoring and performance tracking across distributed services.

Eklavya.me

Jan. 2024 – Apr. 2024

Machine Learning Engineer, Internship, Remote

Goa, India

- Fine-tuned and deployed a **LLaMa-2 7B customer support LLM** on RunPod serverless GPUs using **SGLang**; implemented adaptive RAG with **ChromaDB** over chunked NCERT textbooks, reducing hallucinations and improving response relevance.
- Architected a **multi-agent automated video generation pipeline** (planning → slide content → PPT generation → image synthesis → TTS → video assembly) using **fine-tuned Stable Diffusion** models; reduced educational video production time by **~70%** with human-in-the-loop validation.
- Built distributed media ingestion and delivery pipelines using **Amazon SQS** + **S3**, enabling asynchronous processing of large audio/video assets and improving internal workflow efficiency by **~80%**.

PROJECTS

- [Paper2Pytorch](#): Research paper to code
Tech stack: Next.js 14, FastAPI, Python, Google ADK, Gemini 3 Pro, Server-Sent Events (SSE)
Description:
 - Built a full-stack, multi-agent AI system using Google ADK, FastAPI, and Next.js to convert **ML papers** into executable **PyTorch notebooks**.
 - Designed a 4-stage agent pipeline with Gemini 3 Pro, achieving **~85% zero-shot compilation** via **iterative parsing, design, and auto-repair**.
 - Implemented SSE-based real-time streaming backend and a reactive frontend with GitHub Gists + Colab integration, reducing perceived **latency** by **>40%**.
- [SpeakThought](#) : Audio to Todo tasks
Tech stack: Next.js, Web Audio API (DSP), Google Realtime API, Google ADK
Description:
 - Built a voice-powered task manager using Whisper, GPT-4o, and Google ADK that lets users speak tasks aloud to automatically create todos and send email reminders, achieving **98% extraction accuracy** with a real-time DSP-driven voice visualization reducing perceived **latency** by **40%**.
- [Whisper Mobile](#) : Offline transcription app
Tech stack: Kotlin, Jetpack Compose, Android NDK, JNI, whisper.cpp (ARM64 NEON), Coroutines, StateFlow, MVVM
Description:
 - Designed a 100% offline -Automatic Speech Recognition (ASR) engine using Android NDK and a custom JNI bridge to execute **whisper.cpp C++ inference on ARM64**, achieving **< 4-second latency for 30-second clips**.

EDUCATION

Osmania University

BE, Electronics and Communication Engineering

- 8.41/10 GPA

May 2025

Hyderabad, India

CERTIFICATIONS, SKILLS & ACHIEVEMENTS

- **Certifications:**
 - [Generative AI Leader](#), Google
 - [Blockchain and its applications](#), IIT-K
 - [SQL Advanced](#), HackerRank
- **Achievements:**
 - [First Prize](#), MLH Hackathon, 2023
 - [Second Prize](#), Engenius ideathon, 2024
- **Skills:**
 - Generative AI & LLM Systems:** Multi-Agent Architectures (ADK, LangGraph), Fine-Tuning, Quantization
 - Frontend Development:** React, Angular, TypeScript, HTML, CSS
 - Backend & APIs:** Python, FastAPI, REST APIs
 - Cloud & Distributed Systems:** GCP, AWS (S3, SQS), Docker, Kubernetes, Kafka, RabbitMQ
 - ML & Deployment:** PyTorch, Computer Vision, Vector DBs (ChromaDB, FAISS), Prometheus, Grafana